

# Module: Kernel

Privilege Escalation

Yan Shoshitaishvili  
Arizona State University

# Kernel Memory Corruption

Recall:

```
copy_to_user(userspace_address, kernel_address, length);  
copy_from_user(kernel_address, userspace_address, length);
```

Kernel memory must be kept uncorrupted! Corruption can:

- crash the system
- **brick** the system
- escalate process privileges
- interfere with other processes

All user data should be carefully handled (haha) and only accessed with `copy_to_user` and `copy_from_user`.

# Kernel Vulnerabilities Happen

Kernel code is just code!

Memory corruptions, allocator misuse, etc, all happen in the kernel!

What can you do with this?

# Kernel Race Conditions

Kernel modules are not userspace programs.

- they are always prone to multi-threading
  - what happens if two devices open `/dev/pwn-college` simultaneously?
- they could disappear or swap resources mid-execution
  - what happens if `make_root.ko` is removed while `/proc/pwn-college` is open?

**Race conditions** are huge problems plaguing kernels!

# The Classic: Privilege Escalation

The kernel tracks user the privileges (and other data) of every running process.

```
struct task_struct {
    struct thread_info    thread_info;

    /* -1 unrunnable, 0 runnable, >0 stopped: */
    volatile long         state;

    void                  *stack;
    atomic_t              usage;
    // ...

    int                   prio;
    int                   static_prio;
    int                   normal_prio;
    unsigned int          rt_priority;

    struct sched_info     sched_info;

    struct list_head      tasks;

    pid_t                 pid;
    pid_t                 tgid;

    /* Process credentials: */

    /* Objective and real subjective task credentials (COW): */
    const struct cred __rcu *real_cred;

    /* Effective (overridable) subjective task credentials (COW): */
    const struct cred __rcu *cred;
    // ...
};
```

```
struct cred {
    atomic_t    usage;
    kuid_t      uid;          /* real UID of the task */
    kgid_t      gid;          /* real GID of the task */
    kuid_t      suid;         /* saved UID of the task */
    kgid_t      sgid;         /* saved GID of the task */
    kuid_t      euid;          /* effective UID of the task */
    kgid_t      egid;         /* effective GID of the task */
    kuid_t      fsuid;         /* UID for VFS ops */
    kgid_t      fsgid;         /* GID for VFS ops */
    unsigned    securebits;    /* SUID-less security management */
    kernel_cap_t cap_inheritable; /* caps our children can inherit */
    kernel_cap_t cap_permitted;  /* caps we're permitted */
    kernel_cap_t cap_effective;  /* caps we can actually use */
    kernel_cap_t cap_bset;       /* capability bounding set */
    kernel_cap_t cap_ambient;    /* Ambient capability set */
    // ...
};
```

# How do we set these?

The credentials are supposed to be immutable (i.e., they can be cached elsewhere, and shouldn't be updated in place). Instead, they can be replaced:

```
commit_creds(struct cred *)
```

The cred struct seems a bit complex, but the kernel can make us a fresh one!

```
struct cred * prepare_kernel_cred(struct task_struct  
                                *reference_task_struct)
```

Luckily, if we pass NULL to the reference struct, it'll give us a cred struct with root access and full privileges!

# How do we set these?

We have to run:

```
commit_creds(prepare_kernel_cred(0));
```

# Complications

How do we know where `commit_creds` and `prepare_kernel_cred` are in memory?

1. Older kernels (or newer kernels when kASLR is disabled) are mapped at predictable locations.
2. `/proc/kallsym` is an interface for the kernel to give root these addresses.
3. If enabled, gdb support is your friend.
4. Otherwise, it's the exact same problem as userspace ASLR!